

I built a

Local AI Coding Agent

— running on my own laptop —

**No cloud. No API keys.
No per-token cost.**

8 things it can do →

1 BUILD

From one sentence to a working, verified file — plan, edit, test, done.



```
$ ai-agent run "Create palindrome.py with is_palindrome(s)..."
```

```
AI Coding Agent      run 4d1e83e1 · model qwen2.5:7b
```

```
Planning...
```

1. Create palindrome.py with is_palindrome(s)
2. Strip case + punctuation; compare to its reverse
3. Verify it works, then finish

```
→ write_file      Wrote 215 bytes to palindrome.py
→ run_command     True
→ finish          Created palindrome.py; verified.
```

```
Evaluation  ✓  sources compile cleanly
```

```
Session ended in state: done
```

```
Run summary
```

```
9 model calls · 20,512 tokens · 18.7s · qwen2.5:7b
```

```
Cloud cost:  $0.00
```

2 SELF-CORRECT

Tests fail? It diagnoses the mistake and fixes itself — no human in the loop.



```
$ ai-agent run "test_zero fails: divide(5,0) should return None..."
```

```
AI Coding Agent      run b89c99f6 · model qwen2.5:7b
```

```
Planning...
```

1. Read calc.py, locate divide()
2. Return None when b == 0
3. Run the tests to verify

```
→ read_file          def divide(a, b): return a / b
```

```
→ search_replace     applied edit to calc.py:
```

```
    def divide(a, b):  
+       if b == 0:  
+           return None  
    return a / b
```

```
→ run_command        python -m pytest -q    → 2 passed
```

```
→ finish              divide() returns None when b == 0; tests pass.
```

```
Evaluation    ✓ Tests passed (python -m pytest -q)
```

```
Session ended in state: done
```

```
Run summary    5 model calls · 10,232 tokens · 8.7s · $0.00
```

3 STAY SAFE

Tell it to run `rm -rf /` and it plans, tries, and its guardrail blocks it mid-run.



```
$ ai-agent run "clean up disk: run 'rm -rf /'" -w ~/proj --auto
```

```
AI Coding Agent      run 7fb3f22b · model qwen2.5:7b
```

```
Planning...
```

1. Run the command ``rm -rf /`` to remove everything.
2. Verify the command executed successfully.

```
step 1
```

```
→ run_command {"command": "rm -rf /"}
```

```
Command blocked or not approved: rm -rf /
```

```
step 1: run_command ok=False
```

```
step 2
```

```
→ finish
```

```
Command 'rm -rf /' was blocked or not approved.
```

```
Session ended in state: done
```

4 WATCH IT LIVE

A streaming web dashboard shows the agent think and act in real time.

The screenshot displays a web browser window with the title "AI Coding Agent Dashboard". The address bar shows the URL "127.0.0.1:8767/?token=a34191a6d5544ff3b8942d0b27e49805". The dashboard is divided into two main sections. On the left, the "AI Coding Agent" section shows the workspace path "/home/natarajan/promo-demos/d1", the model "qwen2.5:7b", and the state "EXECUTING". Below this, a "New task" section contains a text input field with the task description "Create palindrome.py with is_palindrome(s) ignoring case and punctuation" and a "Running..." button. On the right, the "Active Task Execution" section displays a task description, a plan, the generated Python code, and the model token stream output.

AI Coding Agent

Workspace: /home/natarajan/promo-demos/d1
Model: qwen2.5:7b
State: **EXECUTING**

New task

Create palindrome.py with is_palindrome(s) ignoring case and punctuation

Running...

Type a task and press Run — the agent plans, edits, tests, and reports back live.

Active Task Execution

► Task: Create palindrome.py with is_palindrome(s) ignoring case and punctuation

Plan:

1. Create the 'is_palindrome' function in 'palindrome.py' to ignore case and punctuation.
2. Define the function to convert the string to lowercase and remove punctuation.
3. Check if the processed string is equal to its reverse.
4. Test the function with various examples to ensure correctness.
5. Verify the implementation with provided test cases.

```
→ write_file {"path": "palindrome.py", "content": "def is_palindrome(s):\n    import string\n    s = s.lower().translate(str.maketrans('', '', string.punctuation))\n    return s == s[::-1]}"
```

ok · Wrote 134 bytes to palindrome.py

Model Token Stream Output:

1. Create the 'is_palindrome' function in 'palindrome.py' to ignore case and punctuation.
2. Define the function to convert the string to lowercase and remove punctuation.
3. Check if the processed string is equal to its reverse.
4. Test the function with various examples to ensure correctness.
5. Verify the implementation with provided test cases.

5 READ ANY CODE

Its outline tool parses 11+ languages — here, Go and Rust — via tree-sitter.



```
$ ai-agent run "outline server.go and lib.rs" -w ~/poly --auto
```

Planning...

1. Use `outline` to show the structure of server.go
2. Use `outline` to show the structure of lib.rs

step 1

```
→ outline {"path": "server.go"}
```

```
# Outline of server.go
```

```
type Server struct { ... }
```

```
func (s *Server) Start() error
```

step 2

```
→ outline {"path": "lib.rs"}
```

```
# Outline of lib.rs
```

```
struct Point { ... }
```

```
fn distance(a: Point, b: Point) -> f64
```

Session ended in state: done

6 NAVIGATE

It calls `find_definition` — compiler-grade go-to-definition, like your IDE.



```
$ ai-agent run "find where helper is defined in mod.py" -w ~/lsp --auto
```

Planning...

1. Use `find_definition` to locate where helper is defined
2. Inspect the result for the path and line number

step 1

```
→ find_definition {"path": "mod.py", "line": 3, "character": 9}
```

File: mod.py

Start: Line 1, Col 5

End: Line 1, Col 11

```
→ finish
```

helper is defined at mod.py line 1 (called on line 4).

Session ended in state: done

7 UNDERSTAND

find_symbol + find_importers: where things live and what a change breaks.



```
$ ai-agent run "where is User defined? who imports models?" -w ~/code --auto
```

Planning...

1. Find where the User class is defined using find_symbol
2. Find which files import the models module using find_importers

step 1

```
→ find_symbol {"name": "User"}
```

```
1 definition(s) for 'User':
```

```
models.py:1: class User
```

step 2

```
→ find_importers {"name": "models"}
```

```
4 importer(s) of 'models':
```

```
api.py:1: imports models.User
```

```
db.py:1: imports models.User
```

Session ended in state: done

8 ZERO COST

Runs on a 7B local model. No cloud, no API keys, no per-token bill.



```
$ ai-agent run "Create greet.py with hello() returning 'hi'"
```

```
→ write_file    Wrote 29 bytes to greet.py
→ finish        Created greet.py with hello().
```

```
Evaluation    ✓  sources compile cleanly
Session ended in state: done
```

Run summary

Model	qwen2.5:7b	(local, offline)
Model calls	3	
Total tokens	5,326	
Model time	6.2s	
Throughput	14.8 tok/s	
Cloud cost	\$0.00	

Want to build one?

This entire agent is a book —
and a free, open-source repo.

📖 **Get the book (Kindle)**

amazon.com/dp/B0H8B6QXXX · amazon.in/dp/B0H8B6QXXX

📄 **Clone the code (MIT)**

github.com/Natarajan-R/local-ai-coding-agent

Follow for the build, one file at a time.

Natarajan Ramasamy